

## Q1 Numerical Answer Formatting

0 Points

The last question of this HW might have answers that are decimal numbers. Due to current limitations of Gradescope, your answers must be an exact string match to ours. In order to ensure an exact match, please carefully follow the following formatting for your numerical answers.

- Do not round decimals. None of the answers are infinite decimals, so include full precision (all answers should be less than 5 places after the decimal).
- Do not include any leading or trailing 0s unless they are necessary to show the location of the decimal
- If the number is an integer, do not include a decimal

### Examples:

.1234

-.001

10.4

-10


0

Note: If you use the Python interpreter to do your math, floating point error may lead to inexact decimal numbers. It is probably best to use another calculator, but if you do use Python you may need to adjust its output to get the actual exact answer.

## Q2 Minimax

11 Points

Consider the zero-sum game tree shown below. Triangles that point up, such as at the top node (root), represent choices for the maximizing player; triangles that point down represent choices for the minimizing player. Outcome values for the maximizing player are listed for each leaf node, represented by the values in squares at the bottom of the tree. Assuming both players act optimally, carry out the minimax search algorithm. Enter the values for the letter nodes in the boxes below the tree.



Input Answers Here

**A:**

1

**B:**

2

**C:**

3

**D:**


3

### Q3 Expectiminimax

12 Points

Consider the game tree shown below. As in the previous problem, triangles that point up, such as the top node (root), represent choices for the maximizing player; triangles that point down represent choices

for the minimizing player. The circular nodes represent chance nodes in which each of the possible actions may be taken with equal probability. The square nodes at the bottom represent leaf nodes. Assuming both players act optimally, carry out the expectiminimax search algorithm. Enter the values for the letter nodes in the boxes below the tree.



Input Answers Here

**A:**

15

**B:**

19

**C:**

8

**D:**

24

**E:**

15

**F:**

8


**G:**

15

**Q4** Unknown Leaf Value

12 Points

Consider the following game tree, where one of the leaves has an unknown payoff,  $x$ . Player 1 moves first, and attempts to maximize the value of the game.



Each of the next 3 questions asks you to write a constraint on  $x$  specifying the set of values it can take. Please specify your answer in one of the following forms:

- Write **All** if  $x$  can take on all values

- Write None if  $x$  has no possible values
- Use an inequality in the form  $x < \{\text{value}\}$ ,  $x > \{\text{value}\}$ , or  $\{\text{value1}\} < x < \{\text{value2}\}$  to specify an interval of values. As an example, if you think  $x$  can take on all values larger than 16, you should enter  $x > 16$ .

Note that the answer check is sensitive to ordering, spacing, and capitalization, so do not use spaces in your answers and make sure the form matches the examples above.

#### Q4.1

3 Points

Assume Player 2 is a minimizing agent (and Player 1 knows this). For what values of  $x$  is Player 1 guaranteed to choose Action 1 for their first move?

$x > 8$

#### Q4.2

3 Points

Assume Player 2 chooses actions at random with each action having equal probability (and Player 1 knows this). For what values of  $x$  is Player 1 guaranteed to choose Action 1?

$x > 9$

#### Q4.3

3 Points

Denote the minimax value of the tree as the value of the root when Player 1 is the maximizer and Player 2 is the minimizer. Denote the expectimax value of the tree as the value of the root when Player 1 is the maximizer and Player 2 chooses actions at random (with equal probability). For what values of  $x$  is the minimax value of the tree worth more than the expectimax value of the tree?

None

**Q4.4**

3 Points


Is it possible to have a game, where the minimax value is strictly larger than the expectimax value?

☐ Yes☒ No**Q5 Alpha-Beta Pruning**

13 Points

Consider the game tree shown below. Triangles that point up, such as at the top node (root), represent choices for the maximizing player; triangles that point down represent choices for the minimizing player. Assuming both players act optimally, use alpha-beta pruning to find the value of the root node. The search goes from left to right; when choosing which child to visit first, choose the left-most unvisited child. In the first set of boxes below, enter the values of the labeled nodes. Then, select the leaf nodes that don't get visited due to pruning.

Hint: Note that the value of a node where pruning occurs is not necessarily the maximum or minimum (depending on which node) of its children. When you prune on conditions  $V > \beta$  or  $V < \alpha$ , assume that the value of the node is  $V$ .



Enter the values of the labeled nodes

**A:**

4

**B:**

3

**C:**

2

**D:**

4

Select the boxes for leaf values that don't get visited due to pruning

☐ 8☐ 4☐ 9☐ 3☒ 5☒ 6☐ 7☐ 2☒ 1

## Q6 Non-Zero-Sum Games

8 Points

### Q6.1

4 Points

The standard minimax algorithm calculates worst-case values in a zero-sum two player game, i.e. a game for which in all terminal states  $s$ , the utilities for players A (MAX) and B (MIN) obey


$U_A(s) + U_B(s) = 0$ . In this zero-sum setting, we know that  $U_A(s) = -U_B(s)$ , so we can think of player B as simply minimizing  $U_A$ .

In this problem, you will consider the non-zero-sum generalization, in which the sum of the two players' utilities are not necessarily zero. The leaf utilities are now written as pairs  $(U_A, U_B)$ . In this generalized setting, A seeks to maximize  $U_A$ , the first component, while B seeks to maximize  $U_B$ , the second component.

Consider the non-zero-sum game tree below. Note that left-pointing triangles (such as the root of the tree) correspond to player A, who maximizes the first component of the utility pair, whereas right-pointing triangles (nodes on the second layer) correspond to player B, who maximizes the second component of the utility pair.

Propagate the terminal utility pairs up the tree using the appropriate generalization of the minimax algorithm on this game tree. In case of ties, choose the leftmost child. Select the correct values for the letter nodes below the tree.





Your answer should be in the format  $(X,Y)$ , where  $X$  is the value of Player A and  $Y$  is the value of Player B at a node (for example " $(7, 1)$ "). Note that the answer check is sensitive to formatting, so be sure to include the parenthesis and do not use any spaces in your answers.

**A:**

(2,14)

**B:**

(16,13)

**C:**

(15,11)

**D:**

(16,13)

## Q6.2

4 Points

In this problem, you will again consider the non-zero-sum generalization, in which the sum of the two players' utilities are not necessarily zero. The leaf utilities are now written as pairs  $(U_A, U_B)$ .

In this generalized setting, A seeks to maximize  $U_A$ , the first component, while B seeks to maximize  $U_B$ , the second component.

Assume that your generalization of the minimax algorithm calculates a value  $(U_A^*, U_B^*)$  for the root of the tree. Assume no utility value for A or for B appears more than once in the terminal nodes (this means there will be no need for tie-breaking). Which of the following statements are true?

☒ Assuming A and B both play optimally, player A's outcome is guaranteed to be exactly  $U_A^*$ .

☒ Assuming A and B both play optimally, player B's outcome is guaranteed to be exactly  $U_B^*$ .

☐ Assuming B plays sub-optimally (but A plays optimally), A's outcome is guaranteed to be at least  $U_A^*$ .


## Q7 Possible Pruning

10 Points

### Q7.1

5 Points

Assume we run  $\alpha - \beta$  pruning, expanding successors from left to right, on a game with tree as shown in Figure (a) below.



Which of the following statements are true?

☒ There exists an assignment of utilities to the terminal nodes such that no pruning will be achieved (shown in Figure (a)).


☒ There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (b) will be achieved.

☐ There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (c) will be achieved.

☐ None of the above.

## Q7.2

5 Points



☐ There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (d) will be achieved.

☐ There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (e) will be achieved.

☐ There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (f) will be achieved.

☒ None of the above.

## Q8 Suboptimal Strategies

9 Points

**Q8.1**

3 Points

Player MAX and player MIN are playing a zero-sum game with a finite number of possible moves. MAX calculates the minimax value of the root to be  $M$ . You may assume that at every turn, each player has at least 2 possible actions. You may also assume that a different sequence of moves will always lead to a different score (i.e., no two terminal nodes have the same score). Which of the following statements are true?

- ☒ Assume MIN is playing sub-optimally at every turn, but MAX does not know this. The outcome of the game could be larger than  $M$  (i.e. at least as good a MAX).
- ☐ Assume MIN is playing sub-optimally at every turn. If MAX plays according to the minimax strategy, the outcome of the game could be less than  $M$ .

**Q8.2**

3 Points

For this question, assume that MIN is playing randomly (with a uniform distribution) at every turn, and MAX knows this. Which of the following statements are true?

- ☐ There exists a policy for MAX such that MAX can guarantee a better outcome than  $M$ .
- ☒ There exists a policy for MAX such that MAX's expected outcome is better than  $M$ .
- ☐ To maximize his or her expected outcome, MAX should play according to the minimax strategy (i.e. the strategy that assumes MIN is playing optimally).

**Q8.3**

3 Points

Which of the following statements are true?

- ✓ Assume MIN is playing sub-optimally at every turn. MAX following the minimax policy will guarantee a better outcome than  $M$ .
- ✓ Assume MIN is playing sub-optimally at every turn, and MAX knows exactly how MIN will play. There exists a policy for MAX to guarantee a better outcome than  $M$ .

**Q9 Shallow Search**

12 Points

In this question, we will investigate shallow search, also known as depth-limited search. Depth-limited search is not guaranteed to find the optimal solution to the original problem. The point of this question is to explore some of the (potentially undesirable) behavior of depth-limited search, and to illustrate that the quality of the evaluation function can play a big role in how well depth-limited search performs.

Consider the following Pacman configuration, in the board below. At each time step, Pacman can move either West (left) or East (right) and is using limited-depth minimax search (where the minimizing agent does not really do anything) to choose his next move. Pacman is 3 East moves away from the food. We will consider the following state evaluation functions:

- $F1(\text{state}) = -(\text{Number of food pellets left})$
- $F2(\text{state}) = -(\text{Number of food pellets left}) + 0.5/(\text{distance to closest food pellet} + 1)$ ; distance to closest food pellet is taken as 0 when no food remains.

The search depth referred to in this question corresponds to the depth in a search tree that only considers the maximizer's actions. For example, if the search considers sequences of up to 2 actions by the maximizer, it'd have a search depth of 2.

Note that there can be ties (in which case both East or West could be returned by the search). Also, note that a search does not finish when

the dots are eaten.



Using F1 as the state evaluation function, for what search depths could East be returned by the search?

☒ 1

☒ 2

☒ 3

☒ 4

☐ None of the above.

Using F1 as the state evaluation function, for what search depths could West be returned by the search?

☒ 1

☒ 2

☐ 3

☐ 4

☐ None of the above.

Using F2 as the state evaluation function, for what search depths could East be returned by the search?

☒ 1☒ 2☒ 3☒ 4☐ None of the above.

Using F2 as the state evaluation function, for what search depths could West be returned by the search?

☐ 1☐ 2☐ 3☐ 4☒ None of the above.

## Q10 Solving MDPs

13 Points

Consider the gridworld MDP for which **Left** and **Right** actions are 100% successful.

Specifically, the available actions in each state are to move to the neighboring grid squares. From state  $a$ , there is also an exit action available, which results in going to the terminal state and collecting a reward of 10. Similarly, in state  $e$ , the reward for the exit action is 1. Exit actions are successful 100% of the time.

10				1
----	--	--	---	---

a

b

c

d

e

Let the discount factor  $\gamma = 1$ . Fill in the following quantities.

$$V_0(d) =$$

0

$$V_1(d) =$$

0

$$V_2(d) =$$

1

$$V_3(d) =$$

1

$$V_4(d) =$$

10

$$V_5(d) =$$

10



## HW 2 (Electronic Component)

● GRADED

## STUDENT

Benjamin Cheung

## TOTAL POINTS

**100 / 100 pts**

## QUESTION 1

Numerical Answer Formatting

**0 / 0 pts**

## QUESTION 2

Minimax

**11 / 11 pts**

## QUESTION 3

Expectiminimax

**12 / 12 pts**

## QUESTION 4

Unknown Leaf Value

**12 / 12 pts**

4.1 (no title)

**3 / 3 pts**

4.2 (no title)

**3 / 3 pts**

4.3 (no title)

**3 / 3 pts**

4.4 (no title)

**3 / 3 pts**

## QUESTION 5

Alpha-Beta Pruning

**13 / 13 pts**

## QUESTION 6

Non-Zero-Sum Games

**8 / 8 pts**

6.1 (no title)

**4 / 4 pts**

6.2 (no title)

**4 / 4 pts**

## QUESTION 7

Possible Pruning

**10 / 10 pts**

7.1 (no title)

**5 / 5 pts**

7.2 (no title)

**5 / 5 pts**

**QUESTION 8**

## Suboptimal Strategies

**9 / 9 pts**

8.1 (no title)

**3 / 3 pts**

8.2 (no title)

**3 / 3 pts**

8.3 (no title)

**3 / 3 pts****QUESTION 9**

## Shallow Search

**12 / 12 pts****QUESTION 10**

## Solving MDPs

**13 / 13 pts**